



eBPF

PB173 Kernel Development Learning Pipeline

Viktor Malík

Principal Software Engineer

What we will cover:

- ▶ General overview
 - Evolution from classic BPF, basic principles, applications
- ▶ eBPF for system observability
 - Typical workflow, attach points, events
- ▶ How to write eBPF programs
 - Available frameworks (`libbpf`, `bpftrace`)
- ▶ Underlying concepts
 - Verifier, maps, helpers/kfuncs
- ▶ BPF Type Format (BTF)
 - Compact debugging information to boost eBPF capabilities



General overview



What is eBPF?

Extended Berkeley Packet Filter

- ▶ **In-kernel virtual machine** allowing to run custom (sandboxed) programs
- ▶ No need to modify kernel source code or load modules
- ▶ Programs are written using **BPF instructions**
 - JIT-compiled to machine instructions
- ▶ Safety is assured by **BPF verifier**



What is eBPF?

Extended Berkeley Packet Filter

- ▶ **In-kernel virtual machine** allowing to run custom (sandboxed) programs
 - ▶ No need to modify kernel source code or load modules
 - ▶ Programs are written using **BPF instructions**
 - JIT-compiled to machine instructions
 - ▶ Safety is assured by **BPF verifier**
-
- ▶ eBPF is one of the most exciting and developed areas of the Linux ecosystem these days
 - 5-10 new commits per day over the last 5 years
 - ▶ *It's like putting JavaScript into the Linux Kernel - Brendan Gregg*



Evolution from classic BPF

A little bit of history

▶ **Berkeley Packet Filter (BPF)**

- Developed in 1990s for fast packet filtering.
- Two registers, few instructions, very limited memory (512B).
- Now referred to as **classic BPF (cBPF)**
- Not used anymore (implemented by eBPF)



Evolution from classic BPF

A little bit of history

▶ **Berkeley Packet Filter (BPF)**

- Developed in 1990s for fast packet filtering.
- Two registers, few instructions, very limited memory (512B).
- Now referred to as **classic BPF (cBPF)**
- Not used anymore (implemented by eBPF)

▶ **Extended BPF (eBPF)**

- Introduced in 2014 into the Linux kernel.
- More registers (11), instructions, more memory, verification.
- Today **BPF = eBPF** (abbreviation no longer translated).

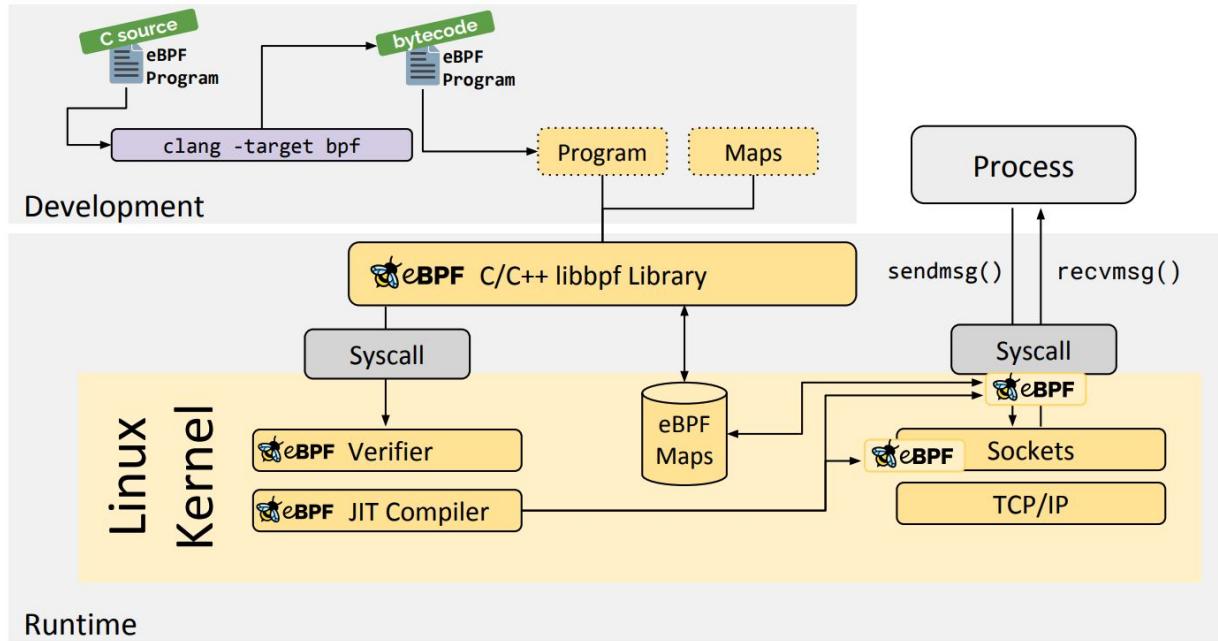


Applications

- ▶ **System observability (tracing)**
 - Attaching eBPF programs to various events in the kernel/userspace
 - **Collecting and processing** information about the events
- ▶ **Networking**
 - Fast packet processing (XDP - eXpress Data Path)
 - The packets can be processed before it reaches the kernel
- ▶ **Security (BPF LSM)**
 - Using BPF programs to implement Mandatory Access Control (MAC) using Linux Security Modules (LSM)
- ▶ **Scheduling (sched_ext)**
 - As of kernel v6.13, it is possible to implement custom schedulers using eBPF
- ▶ **More in development** (e.g. OOM handling)



BPF architecture



eBPF for system observability



Typical workflow of a BPF tracing program

- ▶ Typical workflow:
 - a. eBPF program is **attached to an event** in the system
 - b. When the event fires, the attached BPF program **is executed**
 - c. The BPF program **collects data** about the event and sends it to userspace
 - d. In userspace, data is post-processed, cleaned, and **presented to the user**



Typical workflow of a BPF tracing program

- ▶ Typical workflow:
 - a. eBPF program is **attached to an event** in the system
 - b. When the event fires, the attached BPF program **is executed**
 - c. The BPF program **collects data** about the event and sends it to userspace
 - d. In userspace, data is post-processed, cleaned, and **presented to the user**
- ▶ BPF allows to do a lot of aggregation directly in the BPF program (i.e. kernel)
 - Less data needs to be sent to userspace - **significant performance gain**
 - Important advantage over other observability tools



Available events

- ▶ **Kernel**
 - Static built-in **tracepoints**
 - Dynamic function events
 - **k(ret)probes** - legacy interface, available for any instruction
 - **BPF trampolines** (fentry/fexit) - preferred interface, only available for function boundaries
 - **Memory watchpoints** - an event is generated whenever the observed memory location is accessed (can distinguish read/write/execute access).



Available events

- ▶ **Kernel**
 - Static built-in **tracepoints**
 - Dynamic function events
 - **k(ret)probes** - legacy interface, available for any instruction
 - **BPF trampolines** (fentry/fexit) - preferred interface, only available for function boundaries
 - **Memory watchpoints** - an event is generated whenever the observed memory location is accessed (can distinguish read/write/execute access).
- ▶ **Userspace**
 - Static **USDT** (user-level dynamic tracing) tracepoints
 - Dynamic **uprobes** (analogy of kprobes but for userspace)



How to write eBPF applications



Most popular frameworks

- ▶ **libbpf**
 - **canonical way** of writing BPF programs
 - both kernel and userspace parts written in C
 - many features (CO-RE, global variables) with **first-class support from kernel**
- ▶ BCC (BPF Compiler Collection)
 - userspace part written in Python, BPF part written in C, embedded as a string
 - simpler parsing and presentation of collected data (histograms, etc.)
 - uses libbpf to load programs
- ▶ **bpftrace**
 - custom **high-level language** (similar to SystemTap, DTrace)
 - great for fast prototyping
 - doesn't require deep knowledge of eBPF



libbpf

- ▶ The default **userspace library** for interacting with the kernel
- ▶ A bit complicated to use, especially for newcomers
- ▶ Good starting point: <https://github.com/libbpf/libbpf-bootstrap>
- ▶ Both the BPF and the userspace parts are **written in C**
 - BPF part is compiled into a so-called *BPF skeleton* and embedded in the userspace part
- ▶ Provides a lot of macros and helpers for writing BPF code



libbpf

BPF (kernel) part

- ▶ `read.bpf.c`:

```
SEC("kprobe/vfs_read") ← attach points defined by special macros
int bpf_prog(struct pt_regs *ctx) ← context depends on probe type
{
    __u32 size = (__u32) PT_REGS_PARM3(ctx); ← access to the function argument

    bpf_printk("Reading %d bytes.", size); ← write to special output

    return 0;
}
```



libbpf

Userspace part

- ▶ read.c:

```
#include "read.skel.h" ← include the skeleton (BPF program)
...
int main() {
    struct read_bpf *skel = read_bpf__open();
    int err = read_bpf__load(skel);
    err = read_bpf__attach(skel);
    for (;;) {
        sleep(1);
    }
cleanup:
    read_bpf__destroy(skel);
    return -err;
}
```



bpftrace

- ▶ High-level **tracing language** for Linux using eBPF under the hood
- ▶ A single bpftrace program to create **both the BPF and the userspace** part
- ▶ More information: <https://bpftrace.org/>
- ▶ Allows to write powerful one-liners, great for **fast prototyping**



bpftrace

Examples

- ▶ List all opened files (system wide) by thread name:

```
# bpftrace -e 'tracepoint:syscalls:sys_enter_openat { printf("%s %s\n", comm, str(args.filename)); }'  
Attaching 1 probe...  
ptyxis-agent /proc/1435631/cmdline  
Chrome_IOThread /dev/shm/.org.chromium.Chromium.ogaC8I  
Chrome_IOThread /dev/shm/.org.chromium.Chromium.ogaC8I  
Chrome_IOThread /dev/shm/.org.chromium.Chromium.syZXS4  
mongod /var/lib/mongo/journal  
[...]
```



bpftrace

Examples

- ▶ Show numbers of VFS (virtual filesystem) operations over 1 second:

```
# bpftrace -e 'kprobe:vfs_* { @[func] = count(); } interval:s:1 { exit(); }'  
Attaching 2 probes...
```

```
@[vfs_statx_path]: 9  
@[vfs_statx]: 124  
@[vfs_open]: 197  
@[vfs_fstat]: 240  
@[vfs_getattr_nosec]: 249  
@[vfs_write]: 315  
@[vfs_read]: 1189
```



bpftrace

Examples

- ▶ Read size distribution by thread name:

```
# bpftrace -e 'tracepoint:syscalls:sys_exit_read { @[comm] = hist(args.ret); }'
Attaching 1 probe...
```

[...]@[thunderbird]:

[1] 15 | [@evelynhansen](#)

$$[2, \ 4) \qquad \qquad \theta$$

$$[4, \ 8) \qquad \qquad \qquad \emptyset$$

$$[8, \ 16) \qquad \qquad \qquad \theta$$

[16, 32) 0

[32, 64) 0

[64, 128) 0

[128, 256] 0

[238, 312) 0
[512, 116) 0

[1K, 2K)

[3K–4K) 31

[10]



bpftrace

Prewritten tools

- ▶ bpftrace also comes with a set of **prewritten tools**
- ▶ For example bashreadline.bt:

```
# cat /usr/share/bpftrace/tools/bashreadline.bt
[...]
uretprobe:/bin/bash:readline
{
    time("%H:%M:%S  ");
    printf("%-6d %s\n", pid, str(retval));
}

# /usr/share/bpftrace/tools/bashreadline.bt
Attaching 2 probes...
Tracing bash commands... Hit Ctrl-C to end.
TIME      PID      COMMAND
10:06:59  1436892  cat /etc/passwd
```



bpftool

- ▶ The reference CLI tool for **inspection** and **management** of eBPF objects
- ▶ Uses libbpf under the hood
- ▶ Provides CLI for
 - listing, dumping, loading, attaching **BPF programs**,
 - listing, dumping, creating, manipulating **BPF maps**,
 - generating **BPF skeletons**,
 - inspecting **BPF Type Information (BTF)**,
 - showing **BPF features** available on the system,
 - and much more!



bpftool

Managing eBPF programs

- ▶ List running eBPF programs

```
# bpftool prog list
[...]
44874: cgroup_device  name sd_devices  tag 89ca0032d27da4fb  gpl
    loaded_at 2025-11-10T07:47:16+0100  uid 0
    xlated 1952B  jited 1201B  memlock 4096B
    pids systemd(1)      ← modern systemd loads BPF programs
44917: tracepoint  name tracepoint_syscalls_sys_enter_openat_1  tag 562fe9f18ea5d08f  gpl
    loaded_at 2025-11-10T11:22:28+0100  uid 0
    xlated 1208B  jited 648B  memlock 4096B  map_ids 33395,33396,33391,33392,33394
    btf_id 52803
    pids bpftrace(549407)  ← our bpftrace program
```



bpftool

Managing eBPF programs

- ▶ Show instructions of a BPF program

```
# bpftool prog dump xlated id 44917
int64 tracepoint_syscalls_sys_enter_openat_1(int8 * ctx):
;
    0: (bf) r7 = r1                      ← register assignment
    1: (18) r1 = map[id:33416][0]+0      ← map creation
    3: (79) r8 = *(u64 *)(r1 +0)         ← memory access
    4: (b7) r0 = -1837465548
    5: (bf) r0 = &(void __percpu *)(r0)
    6: (61) r0 = *(u32 *)(r0 +0)
    [...]
   17: (85) call bpf_probe_read_kernel#-123088 ← function (BPF helper) call
   18: (7b) *(u64 *)(r6 +0) = r9
   19: (7b) *(u64 *)(r10 -32) = r9
   20: (7b) *(u64 *)(r10 -40) = r9          ← stack access (r10 is frame pointer)
   [...]
```



bpftool

Managing eBPF maps

- ▶ List created eBPF maps

```
# bpftool map list
[...]
33449: percpu_hash name AT_ flags 0x0
    key 8B  value 8B  max_entries 4096  memlock 722112B  ← key and value types
    btf_id 52913
    pids bpftrace(550454)
[...]
```

- ▶ Dump map contents

```
# bpftool map dump id 33449
[{
    "key": -1887304311,
    "values": [
        {
            "cpu": 0,
            "value": 0
        },
        ...
    ]
[...]
```



Underlying concepts



eBPF Verifier

- ▶ The most important component of the system – **every program must pass** the verifier
- ▶ Properties checked:
 - **Memory access safety**
 - BPF program cannot access memory outside of its context.
 - Pointer dereferencing must use special helper functions.
 - **Stack and register access safety** – no reading uninitialized stack/register values
 - **Instruction reachability** – no dead or unreachable instructions
 - **Termination**
 - eBPF programs must terminate
 - At most 1 million instructions are allowed
 - ...and several others



eBPF maps

- ▶ Generic key-value storage accessible from **both kernel and userspace**
- ▶ Typically used for multiple purposes:
 - **Passing collected data** from BPF programs to userspace
 - **Sharing state** between (instances of) BPF programs

Example with bpftrace:

```
kprobe:vfs_* { @timestamp[tid] = nsecs }  
kretprobe:vfs_* { printf("%s ran for %d ns\n", func, nsecs - @timestamp[tid]); }
```

- ▶ Various map kinds available: (per-cpu) hash maps and arrays, queues, stacks, ...



eBPF helpers and kernel functions

- ▶ Due to safety reasons, eBPF programs **cannot call arbitrary kernel functions**
- ▶ There are lists of functions which are safe (safety is ensured by the verifier) to run:
 - accessing information about the **running process** – name, PID, TID, curtask, stacktrace, ...,
 - map manipulation (e.g. inserting, finding, and deleting elements),
 - **accessing memory**,
 - **iteration**,
 - ... and many many more.
- ▶ These functions are of 2 kinds:
 - BPF **helpers** - stable, legacy, not added anymore.
 - BPF kernel functions (**kfuncs**) – “unstable”, preferred, new are added all the time.



Looping/iteration

General description

- ▶ Problem: eBPF programs **must terminate**, otherwise the system would hang.
- ▶ In the first version, the verifier prohibited loops completely.
- ▶ Then, **bounded loops** were added.
- ▶ Still, some **problems remained**:
 - It is still hard for the verifier to prove that a loop is bounded.
 - It is desirable to iterate over collections which are finite but the number of elements is not known beforehand.



Looping/iteration

Current state

- ▶ These days, there are many options for using loops in BPF programs:

- **bpf_loop helper** allows to execute a function a number of times
 - Simple verification (if the function terminates, the loop terminates)
- **eBPF iterators** are special program types that allow executing code for entry from some collection of kernel objects (running tasks, virtual memory areas, TCP sockets, ...)

```
SEC("iter/task")
int iter_task(struct bpf_iter__task_file *ctx) {
    struct task_struct *task = ctx->task;
    // Do something with the task pointer
}
```

- **Open-coded iterators** allow to iterate these collections from within other BPF programs

```
bpf_for_each(task, task_ptr, NULL, BPF_TASK_ITER_ALL_PROCS) {
    // Do something with the task pointer
}
```



Accessing memory

- ▶ eBPF programs **cannot access arbitrary memory**
 - Possible **out-of-bounds access** (can be checked by the verifier)
 - Possible **page faults**
- ▶ For accessing potentially unsafe memory, special helpers must be used:

```
SEC("kprobe/vfs_read")
int bpf_prog(struct pt_regs *ctx)
{
    struct task_struct *current = (struct task_struct *)bpf_get_current_task();
    struct task_struct *parent;
    bpf_probe_read_kernel(&parent, sizeof(parent), &current->real_parent);
}
```



BPF Type Format (BTF)



BPF Type Format (BTF)

- ▶ Problem: a lot of information useful for tracing is in the debugging information
 - Names of variables, parameters, struct fields, etc.
 - But DWARF is too large (kernel-debuginfo has 4.4 GB on Fedora)
- ▶ BTF is a **compact format for kernel debugging information**
 - Contains definitions of all kernel functions
 - Generated from DWARF by deduplication (4.5 MB BTF vs 195 MB DWARF)
- ▶ Thanks to the small size, BTF is **embedded** in most kernels by default
 - See for yourself (`/sys/kernel/btf/vmlinux`)



Compile-Once, Run-Everywhere (BTF)

Features enabled by BTF

- ▶ Problem: **layout of kernel structures can change** between versions (no stable ABI)
- ▶ Normally, eBPF programs would have to be recompiled for each kernel version
- ▶ CO-RE uses BTF to **dynamically adjust** the BPF program to the current kernel upon loading
- ▶ Example using libbpf:

```
SEC("kprobe/vfs_read")
int bpf_prog(struct pt_regs *ctx)
{
    struct task_struct *current = (struct task_struct *)bpf_get_current_task();
    int ppid = BPF_CORE_READ(current, real_parent, tgid);
}
```



BPF trampolines (fentry/fexit)

Features enabled by BTF

- ▶ New probe type for attaching BPF programs to function entries/exits.
- ▶ Advantages:
 - Practically no overhead (use special nop instructions)
 - Have access to **function arguments by name** (thanks to BTF)
 - **Direct dereferencing** is possible
- ▶ Example using libbpf:

```
SEC("fentry/vfs_open")
int BPF_PROG(vfs_open, const struct path *path, struct file *file)
{
    ... path->dentry->d_name.name ...
}
```



vmlinux.h

Features enabled by BTF

- ▶ When using kernel types, BPF programs need to **include kernel headers**

```
#include <linux/sched.h>
SEC("kprobe/vfs_read")
int bpf_prog(struct pt_regs *ctx)
{
    struct task_struct *current = (struct task_struct *)bpf_get_current_task();
    ... current->tid ...
}
```



vmlinux.h

Features enabled by BTF

- When using kernel types, BPF programs need to **include kernel headers**

```
#include <linux/sched.h>
SEC("kprobe/vfs_read")
int bpf_prog(struct pt_regs *ctx)
{
    struct task_struct *current = (struct task_struct *)bpf_get_current_task();
    ... current->tid ...
}
```

- BTF has all the types so we can use it to generate **vmlinux.h - the header with all kernel types**

```
# bpftool btf dump file /sys/kernel/btf/vmlinux format c > vmlinux.h
```

- Then, all BPF programs just need to include vmlinux.h, nothing else

```
#include "vmlinux.h"
SEC("kprobe/vfs_read")
...
...
```



Conclusion



Conclusion

- ▶ **eBPF is an exciting technology** which is getting a lot of traction these days.
- ▶ One (but not the only one!) of the use-cases is for **system observability**.
- ▶ eBPF completely redefined the way Linux kernel can be extended at runtime.
- ▶ More resources:
 - eBPF website: <https://ebpf.io/>
 - eBPF docs: <https://docs.ebpf.io/>
 - eBPF documentary: <https://ebpfdocumentary.com/>





Thank you

Red Hat is the world's leading provider of enterprise open source software solutions. Award-winning support, training, and consulting services make Red Hat a trusted adviser to the Fortune 500.



[linkedin.com/company/red-hat](https://www.linkedin.com/company/red-hat)



[facebook.com/redhatinc](https://www.facebook.com/redhatinc)



[youtube.com/user/RedHatVideos](https://www.youtube.com/user/RedHatVideos)



twitter.com/RedHat

